

musl security features

MaskRay

<https://maskray.me>

- 1 Introduction
- 2 Miscellaneous
- 3 Stack Smashing Protector
- 4 Secure mode
- 5 malloc
- 6 mallocng data structures
- 7 glibc
- 8 Refereces

musl

- A libc, an implementation of the user-space side of standard C/POSIX functions with Linux extensions.
- General-purpose, not specific to the embedded domain
- Launched in 2011. Milestone 1.0 released in 2014.

Distributions using musl

- Alpine Linux, Dragora, sabotage, Talos, Void Linux
- Shipping musl as an optional package: Gentoo Linux, OpenWrt

Core Principles

- Extracted from Rich Felker's talk *Transitioning from uclibc to musl for embedded development*
- Simplicity as the core approach to size, performance, security, and maintainability
- Factoring for minimal code duplication
- Ease of navigating and understanding code
- Robustness/fail-safety
- Not depending on fancy compiler/toolchain features
- First-class status for UTF-8, non-ASCII characters

Miscellaneous

- Eager/deferred binding instead of lazy binding [wiki](#)
- v1.0.0: PT_GNU_RELRO: some sections are readonly after relocation resolving. `mprotect` such sections.
- v1.1.10: static pie (earlier than [glibc](#))
- v1.1.20: `explicit_bzero`: `bzero` with a compiler barrier
- v1.1.24: `secure_getenv`: `getenv` or (set-user-ID or set-group-ID) no-op
- v1.2.2: `reallocarray`: `realloc` with multiplication overflow check (OpenBSD 5.6)

Stack Smashing Protector

- `-fstack-protector{-explicit,, -string, -all}`
- Add a secret value (canary) after local variables (before the return address) on the stack.
- On return, check whether the canary stays the same.
- `-mstack-protector-guard=` decides whether the secret is stored.

```
void foo(const char *a, const char *b) {
    char buf[16];
    strcpy(buf, a); strcat(buf, b); puts(buf);
}
```

```
foo:                                # @foo
    # Save callee-saved registers.
    pushq %r14; pushq %rbx
    # Allocate local variables and canary.
    subq $24, %rsp
    # Copy arguments to callee-saved registers.
    movq %rsi, %r14; movq %rdi, %rsi
    # Set canary.
    movq %fs:40, %rax # -mstack-protector-guard=tls
    movq __stack_chk_guard(%rip), %rax # -mstack-protector-guard=global
    movq %rax, 16(%rsp)
    movq %rsp, %rbx
    # Main body.
    movq %rbx, %rdi; callq strcpy
    movq %rbx, %rdi; movq %r14, %rsi; callq strcat
    movq %rbx, %rdi; callq puts
    # Check canary.
    movq %fs:40, %rax # -mstack-protector-guard=tls
    movq __stack_chk_guard(%rip), %rax # -mstack-protector-guard=global
    cmpq 16(%rsp), %rax; jne .LBB0_2
    # Epilogue
    addq $24, %rsp; popq %rbx; popq %r14
    retq
.LBB0_2:
    callq __stack_chk_fail
.Lfunc_end0:
```


Stack Smashing Protector in libc

- libc has to initialize (tls) `%fs:40` or (global) `__stack_chk_guard`.
- In glibc, only one is supported (arch-specific configure-time decision). [PR glibc/26817](#) (tls may be less secure due to lack of guard page before static TLS block)
- musl supports `-mstack-protector-guard=tls` and `-mstack-protector-guard=global` at the same time.
- v1.1.9: Allow libc itself to be built with `-fstack-protector` (earlier than glibc)

Secure mode

- set-user-ID and set-group-ID (`chmod u+s` and `chmod g+s`)
- ld.so: disallow `LD_LIBRARY_PATH`/`LD_PRELOAD`, non-absolute `$ORIGIN`, obtaining `$ORIGIN` from main executable `/proc/self/exe`
- date/time functions: disallow TZ file other than `/etc/localtime` (TZ=:/usr/local/etc/localtime)
- catopen: disallow `NLSPATH`
- Ensure fd 0/1/2 are open: a badly written set-UID program may open files (occupying fd 0/1/2) and clobber these files when writing to stdout/stderr.

mallocng

- v1.2.1: “Strong hardening against memory usage errors by the caller, including detection of overflows, double-free, and use-after-free, and does not admit corruption of allocator state via these errors.”
- Mixed in-band and out-of-band metadata. Sensitive states are in the out-of-band metadata.
- Allocation granularity is 16 bytes.

Corruption of allocator state

- glibc (as of 2.31): `free()`: double free detected in tcache
2 double free or corruption (fasttop)
- musl: crash without diagnostic

malloc implementations

- dlmalloc (Doug Lea), ptmalloc (pthread malloc), glibc, jemalloc, tcmalloc, OpenBSD omalloc, Chrome PartitionAlloc
- dlmalloc family: metadata (fd/bk/prev_size/size) in a free chunk overlay user data in an in-use chunk

Characteristics

- malloc after free: first-fit effects, similar to glibc (LIFO in tcache/fastbin; FIFO in unsorted bin). However, randomness can be trivially implemented by changing the current least significant bit heuristic. For a request of n , if the slot has $n + UNIT * (k - 1)$ bytes, `enframe` can cycle the allocation in k places, i.e. a slot can be reused without handling an identical resource identifier. This property is more useful for detecting double-free bugs than hardening.
- Main metadata is stored separately. Out-of-bounds write forges the header of the next slot, but it can hardly do harm.
- double free/invalid free: guaranteed crash due to rigorous verification. In glibc, double free of non-top chunk in fastbin cannot be detected (double free or corruption (fasttop), other bins have more checks).
- No `malloc_set_zero_contents/malloc_set_pattern_fill_contents` (scudo). Data leak from freed memory is possible.
- No `__malloc_hook` or `__free_hook`.

```
// Singleton: ctx
struct malloc_context {
    uint64_t secret;
#ifdef PAGE_SIZE
    size_t pagesize;
#endif
    // When alloc_meta is invoked for the first time, initialize secret and pagesize.
    int init_done;
    unsigned mmap_counter;
    struct meta *free_meta_head;
    // [avail_meta, avail_meta_count) are available meta objs.
    struct meta *avail_meta;
    // When a new page is allocates, (4096-sizeof(meta_area))/sizeof(meta) meta objs are newly available.
    size_t avail_meta_count, avail_meta_area_count, meta_alloc_shift;
    struct meta_area *meta_area_head, *meta_area_tail;
    // [avail_meta_areas, avail_meta_areas+avail_meta_area_count) are available meta areas.
    unsigned char *avail_meta_areas;
    // Doubly linked meta list by size class,
    struct meta *active[48];
    size_t usage_by_class[48];
    uint8_t unmap_seq[32], bounces[32];
    uint8_t seq;
    // [initial brk(0)+pagesize, brk) stores available meta areas.
    uintptr_t brk;
};
```

A group contains a header (of minimum metadata) and $1 \sim 32$ slots for allocation. Each group is paired with a meta object which references back to the group.

The slot size (stride) is decided by the size class. The number of slots is decided by the size class and its usage.

```

struct meta_area {
    uint64_t check; // ctx.secret
    struct meta_area *next;
    int nslots; // (4096-sizeof(meta_area))/sizeof(meta)
    struct meta slots[];
};

struct meta {
    struct meta *prev, *next;
    struct group *mem;
    // Bitmask of unused slots and bitmask of freed slots. They have no intersection.
    volatile int avail_mask, freed_mask;
    // The index of the last slot, i.e. the number of slots minus 1.
    uintptr_t last_idx:5;
    uintptr_t freeable:1;
    uintptr_t sizeclass:6;
    uintptr_t maplen:8*sizeof(uintptr_t)-12;
};

struct group {
    struct meta *meta;
    unsigned char active_idx:5;
    char pad[UNIT - sizeof(struct meta *) - 1];
    // N slots. A slot starts at offset (1+stride*slot_index)*UNIT.
    unsigned char storage[];
};

```


Each chunk is user data preceded by a 4-byte header ($IB = 4$). For a requested size n , the chunk size is $n + 4$. The slot size needs to be at least as large as $n + 4$.

The canonical placement starts the user data at the slot start (offset: 0) and makes its header overlap the last 4 bytes of the previous slot. (The last 4 bytes of the current slot is reserved by the next chunk). This placement can be used for an unaligned allocation.

A chunk may have some trailing spare bytes, named *reserved*. If *reserved* ≥ 5 , use the last 5 bytes to store a zero canary byte and *reserved*, verified by `get_nominal_size` (called by `malloc` and `free`). This can detect out-of-bounds writes.

```
struct chunk {
    uint8_t zero;
    uint8_t idx; // slot_index | (min(reserved,5)<<5); if nested in a larger chunk, slot_index | 6<<5
    uint16_t offset; // (p-g->mem->storage) / UNIT
    // If reserved bytes >= 5, end[-5] is 0 and *(uint32_t*)(end-4) stores reserved.
    char p[stride-IB];
};

struct chunk32 {
    uint32_t offset; // (p-g->mem->storage) / UNIT
    uint8_t non_zero;
    uint8_t idx; // slot_index | (min(reserved,5)<<5)
    uint16_t zero;
    // If reserved bytes >= 5, end[-5] is 0 and *(uint32_t*)(end-4) stores reserved.
    char p[stride-IB];
};
```

Features available in glibc but unavailable in musl

- Memory protection keys (pkey_*)
- AArch64 Branch Target Identification and Pointer Authentication (assembly annotation and ld.so support (mprotext PROG_BTI))
- AArch64 Memory Tagging Extension (including heap tagging in its malloc implementation)
- x86 IBT and SHSTK (assembly annotation and ld.so support)
- Defense in depth: setjmp/longjmp/__cxa_atexit address mangling

References

- Transitioning From uclibc to musl for Embedded Development, Rich Felker
- <https://dustri.org/b/security-features-of-musl.html>, Julien Voisin